

## Chapter 2: Operating-System Structures



## Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- Operating System Debugging
- Operating System Generation
- System Boot



## Objectives

- To describe the services an operating system provides to users, processes, and other systems
- To discuss the various ways of structuring an operating system
- To explain how operating systems are installed and customized and how they boot




## Operating System Services

- One set of operating-system services provides functions that are helpful to the user:
  - User interface - Almost all operating systems have a user interface (UI)
    - Varies between [Command-Line \(CLI\)](#), [Graphics User Interface \(GUI\)](#), [Batch](#)
  - Program execution - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
  - I/O operations - A running program may require I/O, which may involve a file or an I/O device
  - File-system manipulation - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.



## A View of Operating System Services



## Operating System Services (Cont)

- One set of operating-system services provides functions that are helpful to the user (Cont):
  - Communications – Processes may exchange information, on the same computer or between computers over a network
    - Communications may be via shared memory or through message passing (packets moved by the OS)
  - Error detection – OS needs to be constantly aware of possible errors
    - May occur in the CPU and memory hardware, in I/O devices, in user program
    - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
    - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system





## Operating System Services (Cont)

- Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
  - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
    - Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code
  - **Accounting** - To keep track of which users use how much and what kinds of computer resources
  - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
    - **Protection** involves ensuring that all access to system resources is controlled
    - **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
    - If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.



## User Operating System Interface - CLI

Command Line Interface (CLI) or **command interpreter** allows direct command entry

- Sometimes implemented in kernel, sometimes by systems program
- Sometimes multiple flavors implemented - **shells**
- Primarily fetches a command from user and executes it
  - Sometimes commands built-in, sometimes just names of programs
  - If the latter, adding new features doesn't require shell modification



## User Operating System Interface - GUI

- User-friendly **desktop** metaphor interface
  - Usually mouse, keyboard, and monitor
  - **Icons** represent files, programs, actions, etc
  - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
  - Invented at Xerox PARC
- Many systems now include both CLI and GUI interfaces
  - Microsoft Windows is GUI with CLI "command" shell
  - Apple Mac OS X as "Aqua" GUI interface with UNIX kernel underneath and shells available
  - Solaris is CLI with optional GUI interfaces (Java Desktop, KDE)



## Bourne Shell Command Interpreter

```

$ cat /etc/passwd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
sys:x:3:3:sys:/dev:/usr/sbin/nologin
cron:x:4:4:cron:/var/spool/cron/root:/bin/crontab
ftp:x:5:0:ftp:/var/fsp:/usr/sbin/nologin
operator:x:10:10:operator:/var/spool/cron/root:/bin/crontab
nobody:x:65534:65534:nobody:/nonexistent:/usr/sbin/nologin

```



## The Mac OS X GUI



## System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call use
- Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)
- Why use APIs rather than system calls?

(Note that the system-call names used throughout this text are generic)



## Example of System Calls

- System call sequence to copy the contents of one file to another file

**Example System Call Sequence**

- Acquire input file name
- Write prompt to screen
- Acquire output file name
- Write prompt to screen
- Open the input file
- Create output file
- Loop
  - Read from input file
  - Write to output file
- Until read fails
- Close output file
- Write completion message to screen
- Terminate normally

Operating System Concepts – 8th Edition 2.13 Silberschatz, Galvin and Gagne ©2009

## Example of Standard API

- Consider the ReadFile() function in the Win32 API—a function for reading from a file

```

  return value
  ↓
  BOOL ReadFile (HANDLE file, LPVOID buffer, DWORD bytesToRead, LPDWORD bytesRead, LPOVERLAPPED overlapped);
  ↑
  function name
  
```

- A description of the parameters passed to ReadFile()
  - HANDLE file—the file to be read
  - LPVOID buffer—a buffer where the data will be read into and written from
  - DWORD bytesToRead—the number of bytes to be read into the buffer
  - LPDWORD bytesRead—the number of bytes read during the last read
  - LPOVERLAPPED overlapped—indicates if overlapped I/O is being used

Operating System Concepts – 8th Edition 2.14 Silberschatz, Galvin and Gagne ©2009

## System Call Implementation

- Typically, a number associated with each system call
  - System-call interface maintains a table indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
  - Just needs to obey API and understand what OS will do as a result call
  - Most details of OS interface hidden from programmer by API
    - Managed by run-time support library (set of functions built into libraries included with compiler)

Operating System Concepts – 8th Edition 2.15 Silberschatz, Galvin and Gagne ©2009

## API – System Call – OS Relationship

Operating System Concepts – 8th Edition 2.16 Silberschatz, Galvin and Gagne ©2009

## Standard C Library Example

- C program invoking printf() library call, which calls write() system call

```

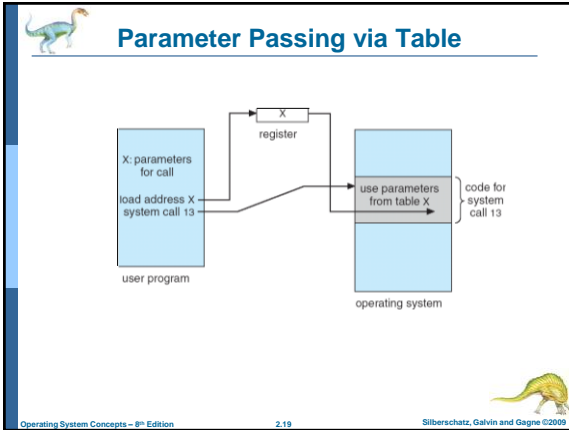
  #include <stdio.h>
  int main ()
  {
    printf ("Greetings");
    return 0;
  }
  
```

Operating System Concepts – 8th Edition 2.17 Silberschatz, Galvin and Gagne ©2009

## System Call Parameter Passing

- Often, more information is required than simply identity of desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
  - Simplest: pass the parameters in registers
    - In some cases, may be more parameters than registers
  - Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
    - This approach taken by Linux and Solaris
  - Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
  - Block and stack methods do not limit the number or length of parameters being passed

Operating System Concepts – 8th Edition 2.18 Silberschatz, Galvin and Gagne ©2009

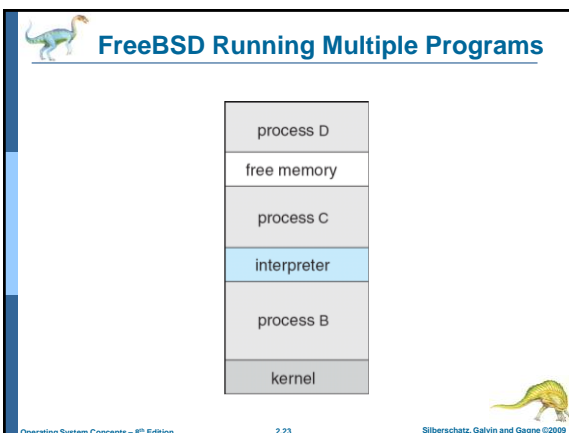
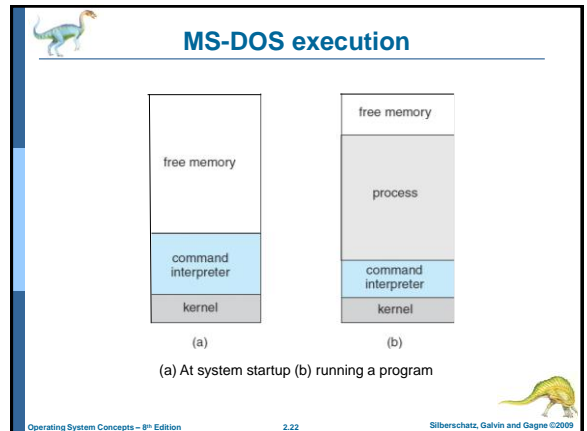


- ### Types of System Calls
- Process control
  - File management
  - Device management
  - Information maintenance
  - Communications
  - Protection
- Operating System Concepts – 8th Edition 2.20 Silberschatz, Galvin and Gagne ©2009

### Examples of Windows and Unix System Calls

Blue box from page 51

Operating System Concepts – 8th Edition 2.21 Silberschatz, Galvin and Gagne ©2009



- ### System Programs
- System programs provide a convenient environment for program development and execution. The can be divided into:
    - File manipulation
    - Status information
    - File modification
    - Programming language support
    - Program loading and execution
    - Communications
    - Application programs
  - Most users' view of the operation system is defined by system programs, not the actual system calls
- Operating System Concepts – 8th Edition 2.24 Silberschatz, Galvin and Gagne ©2009



## System Programs

- Provide a convenient environment for program development and execution
  - Some of them are simply user interfaces to system calls; others are considerably more complex
- File management - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- Status information
  - Some ask the system for info - date, time, amount of available memory, disk space, number of users
  - Others provide detailed performance, logging, and debugging information
  - Typically, these programs format and print the output to the terminal or other output devices
  - Some systems implement a registry - used to store and retrieve configuration information



## System Programs (cont'd)

- File modification
  - Text editors to create and modify files
  - Special commands to search contents of files or perform transformations of the text
- Programming-language support - Compilers, assemblers, debuggers and interpreters sometimes provided
- Program loading and execution- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language
- Communications - Provide the mechanism for creating virtual connections among processes, users, and computer systems
  - Allow users to send messages to one another's screens, browse web pages, send electronic-mail messages, log in remotely, transfer files from one machine to another



## Operating System Design and Implementation

- Design and Implementation of OS not "solvable", but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start by defining goals and specifications
- Affected by choice of hardware, type of system
- User goals and System goals
  - User goals – operating system should be convenient to use, easy to learn, reliable, safe, and fast
  - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient



## Operating System Design and Implementation (Cont)

- Important principle to separate
  - Policy:** What will be done?
  - Mechanism:** How to do it?
- Mechanisms determine how to do something, policies decide what will be done
  - The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later




## Simple Structure

- MS-DOS – written to provide the most functionality in the least space
  - Not divided into modules
  - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated



## MS-DOS Layer Structure



## Layered Approach

- The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers

## Traditional UNIX System Structure

The diagram illustrates the layers of a traditional UNIX system structure:

- (the users)**: shells and commands, compilers and interpreters, system libraries.
- system-call interface to the kernel**: A horizontal line separating user space from kernel space.
- Kernel**: A bracketed section containing:
  - signals terminal handling, character I/O system, terminal drivers
  - file system, swapping block I/O system, disk and tape drivers
  - CPU scheduling, page replacement, demand paging, virtual memory
- kernel interface to the hardware**: A horizontal line separating the kernel from hardware.
- Hardware**: terminal controllers terminals, device controllers disks and tapes, memory controllers physical memory.

## UNIX

- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
  - Systems programs
  - The kernel
    - Consists of everything below the system-call interface and above the physical hardware
    - Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

## Layered Operating System

The diagram shows a layered operating system as concentric circles:

- layer N user interface**: The outermost layer.
- layer 1**: An intermediate layer.
- layer 0 hardware**: The innermost layer, representing the physical hardware.

## Microkernel System Structure

- Moves as much from the kernel into "user" space
- Communication takes place between user modules using message passing
- Benefits:
  - Easier to extend a microkernel
  - Easier to port the operating system to new architectures
  - More reliable (less code is running in kernel mode)
  - More secure
- Detriments:
  - Performance overhead of user space to kernel space communication

## Mac OS X Structure

The diagram shows the Mac OS X structure:

- application environments and common services**: The top layer.
- kernel environment**: A bracketed section containing:
  - BSD**: The user-space kernel.
  - Mach**: The microkernel.

Arrows indicate communication between the application environments and both the BSD and Mach components.

## Modules

- Most modern operating systems implement kernel modules
  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
  - Each is loadable as needed within the kernel
- Overall, similar to layers but with more flexible

Operating System Concepts – 8th Edition 2.37 Silberschatz, Galvin and Gagne ©2009

## Solaris Modular Approach

```

graph TD
    CSK((core Solaris kernel)) --- SC(scheduling classes)
    CSK --- FS(file systems)
    CSK --- LSC(loadable system calls)
    CSK --- EF(executable formats)
    CSK --- SM(STREAMS modules)
    CSK --- MM(miscellaneous modules)
    CSK --- DBD(device and bus drivers)
  
```

Operating System Concepts – 8th Edition 2.38 Silberschatz, Galvin and Gagne ©2009

## Virtual Machines

- A **virtual machine** takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware
- A virtual machine provides an interface *identical* to the underlying bare hardware
- The operating system **host** creates the illusion that a process has its own processor and (virtual) memory
- Each **guest** provided with a (virtual) copy of underlying computer

Operating System Concepts – 8th Edition 2.39 Silberschatz, Galvin and Gagne ©2009

## Virtual Machines History and Benefits

- First appeared commercially in IBM mainframes in 1972
- Fundamentally, multiple execution environments (different operating systems) can share the same hardware
- Protect from each other
- Some sharing of file can be permitted, controlled
- Commute with each other, other physical systems via networking
- Useful for development, testing
- Consolidation** of many low-resource use systems onto fewer busier systems
- "Open Virtual Machine Format", standard format of virtual machines, allows a VM to run within many different virtual machine (host) platforms

Operating System Concepts – 8th Edition 2.40 Silberschatz, Galvin and Gagne ©2009

## Virtual Machines (Cont)

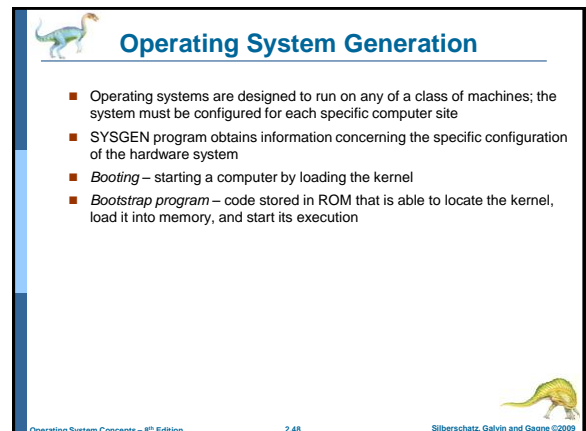
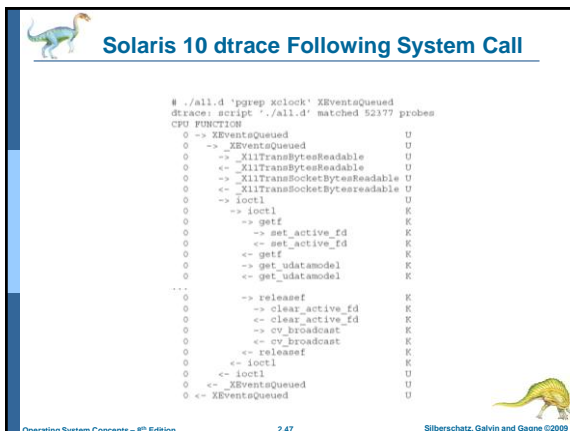
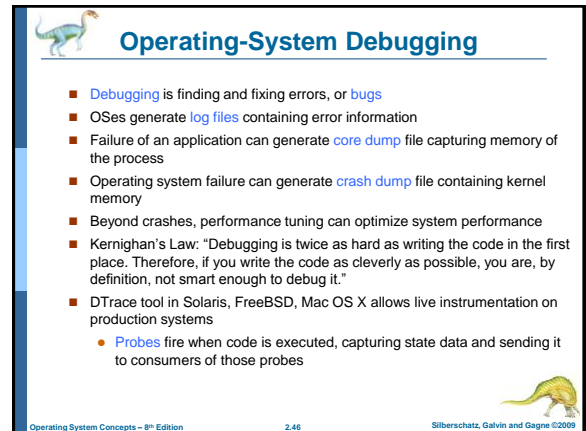
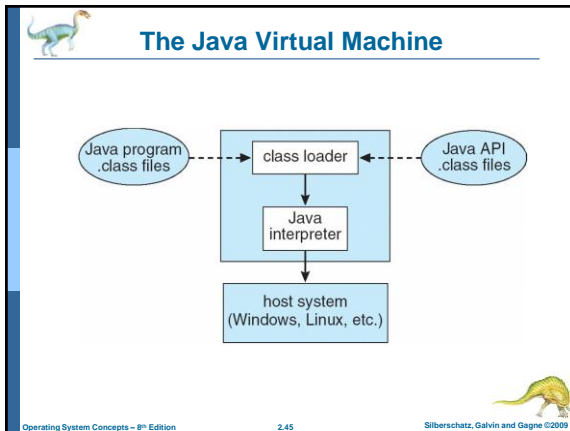
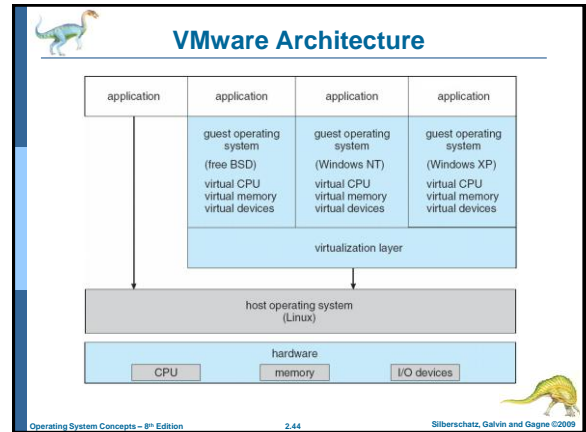
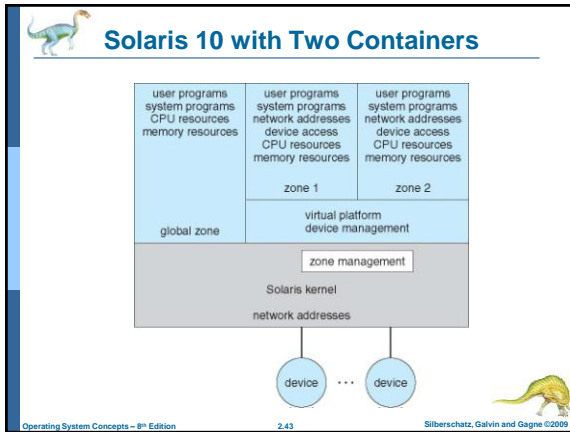
(a) Nonvirtual machine (b) virtual machine

Operating System Concepts – 8th Edition 2.41 Silberschatz, Galvin and Gagne ©2009

## Para-virtualization

- Presents guest with system similar but not identical to hardware
- Guest must be modified to run on paravirtualized hardwareF
- Guest can be an OS, or in the case of Solaris 10 applications running in **containers**

Operating System Concepts – 8th Edition 2.42 Silberschatz, Galvin and Gagne ©2009







## System Boot

- Operating system must be made available to hardware so hardware can start it
  - Small piece of code – **bootstrap loader**, locates the kernel, loads it into memory, and starts it
  - Sometimes two-step process where **boot block** at fixed location loads bootstrap loader
  - When power initialized on system, execution starts at a fixed memory location
    - ▶ Firmware used to hold initial boot code



## End of Chapter 2

